

Why Memory Became the Bottleneck Behind the AI Trade

Everyone bet on the AI **chips** (GPUs, Nvidia). The quieter, bigger story is the AI **memory** that feeds them: made by only three companies, and built in a way that starves the ordinary memory in your phone and laptop. Step by step, this is how it works.

CONTENTS

- 01 Start with the chef and the kitchen
- 02 The counter got too small: the memory wall
- 03 HBM: a giant counter wrapped around the chef
- 04 The sneaky part: the fancy counter steals slabs
- 05 The whole world buys from just three stores
- 06 Why investors care: a toll booth on AI
- 07 What the boss said about 2030
- 08 The catch every trader has to respect
- G Glossary

01 Start with the chef and the kitchen

Picture a **super fast chef** cooking. The kitchen has three parts:

- The **chef is the GPU** (graphics processing unit: the chip that does the work, very fast). It cooks; it does not store.
- The **counter is memory** (the RAM or DRAM chips that hold data while it is being used): small, but right next to the chef.
- The **pantry is storage**: huge, but far away. The chef can only cook with what is already on the counter.

02 The counter got too small: the memory wall

The chef got **so fast** that the counter cannot be restocked in time.

- The chef stands with **empty hands**, waiting. A faster chef just waits more.
- That gap is the **memory wall** (where memory, not the chef, holds dinner up). It is now the real bottleneck of AI.

03 HBM: a giant counter wrapped around the chef

The fix is a special fast counter called **HBM** (high bandwidth memory: the fast counter AI chips need).

- The shelves stack high right next to the chef, not scattered around the kitchen.
- Tiny lifts, like dumbwaiters, called **TSVs** (through silicon vias: wires running straight up through the stacked chips) rush data up.
- Result: a huge counter inches away. **Every serious AI chip needs HBM**, and there is not enough.

04 The sneaky part: the fancy counter steals slabs

Every counter is cut from the same limited slabs of silicon called **wafers** (a factory can only make so many per month).

THE ONE IDEA TO REMEMBER

HBM uses about 2 to 3 times the slab of a plain counter. Every slab used for HBM is **taken away** from the plain memory in phones, laptops, and servers. HBM demand makes **everything** scarce at once, and prices climb on both sides. That double squeeze is the engine.

05 The whole world buys from just three stores

Almost all memory comes from **three companies**, which hands them **pricing power** (raise prices without losing customers). A market like that is an **oligopoly**.

SK Hynix The **HBM leader**, Nvidia's supplier. Filing to list in the US.

Samsung The giant; memory is one slice. Racing to catch up on HBM.

Micron (MU) The **only US listed pure play**. The easy way to own this.

All three are sold out ahead: together they can meet only about **half** of what the big cloud buyers, the **hyperscalers** (Amazon, Microsoft, Google, Meta), have already ordered.

06 Why investors care: a toll booth on AI

A scarce part that **every** AI build needs, sold by only three firms, works as a **toll booth on the whole AI boom**. The **gross margin** (what is left of each sales dollar after making the product; high means real pricing power) shows it.

~70%

of memory chips go to AI data centers in 2026 (est.)

3

firms make almost all the world's DRAM and HBM

~90%

est. HBM gross margin at SK Hynix, above Nvidia's

07 What the boss said about 2030

"AI actually wants to have a lot of HBM, and once you make the HBM, we have to use a lot of wafers."

CHEY TAE-WON · SK GROUP CHAIRMAN · COMPUTEX 2026

SK Hynix guides HBM demand growing about 30% a year through 2030, and warns the shortage could run past 2027. This is a runway of years, not one quarter.

08 The catch every trader has to respect

The honest part: memory is the **most cyclical** (boom and bust) part of chips.

- Like the hottest Christmas toy: everyone wants it, price rockets, factories overbuild, then it **crashes**. Memory has done this for decades.
- **Bulls**: this time demand is real (AI) and supply is disciplined (three players). This cycle is different.
- **Bears**: it is still a commodity; **capex** (spending on new factories) in 2026 to 2027 will overbuild.

So trade it as a **cyclical position, not buy and forget**:

- Set your **invalidation** (the exit signal: prices rolling over, oversupply headlines, inventory builds) before you enter.
- The traders who lose money marry the story at the top. **Plan the exit before the entry.**

THE WHOLE THING IN ONE LINE

The trade is not “AI will be big.” It is: AI cannot run without memory, only three firms make it, and making the fast kind starves the cheap kind.

This is a long position in a structural shortage, with a runway to 2030, until the cycle turns. Know that going in.

SOURCES & NOTE

Sources: SK Hynix HBM demand guidance; Chey Tae-won (SK Group Chairman), Computex 2026; industry estimates for AI memory demand and gross margins.

Educational, not financial advice. Figures marked “est.” are estimates.

Glossary

GPU **graphics processing unit** The chip that does the AI work, fast. The chef who can only cook with what is on the counter.

RAM / DRAM **dynamic random access memory** The counter next to the chef that holds data while cooking: small but close. Storage (the pantry) is huge but far.

Memory wall Where the counter, not the chef, holds dinner up: the chef waiting for ingredients.

HBM **high bandwidth memory** Memory stacked high right next to the chip, restocked at firehose speed. What every AI accelerator needs.

TSV **through silicon via** The tiny vertical lifts (dumbwaiters) that carry data up through the stacked memory.

Wafer The round slab of silicon chips are cut from. A factory's output is roughly fixed: the source of the squeeze.

Oligopoly A market only a few sellers control, which hands them pricing power.

Pricing power The ability to raise prices without losing customers. The whole thesis rests here.

Hyperscaler A giant cloud or AI buyer (Amazon, Microsoft, Google, Meta) that orders memory years ahead.

Gross margin What is left of each sales dollar after making the product. High means strong pricing power.

Pure play A company that does essentially one thing (Micron is roughly only memory), so its stock tracks that market.

Capex **capital expenditure** Money spent building new factories. Too much of it causes the glut that ends the cycle.

Cyclical Prone to boom and bust swings. Memory is the most cyclical part of semiconductors.

Invalidation The preset signal that proves a trade wrong and triggers the exit. Set it before you enter.